

Jiaqi Duan

Fremont, CA | +1 (408) 438-7247 | jd.victoria.work@gmail.com

linkedin.com/in/jiaqi-duan | github.com/Victoriakaey | victoria-duan.vercel.app

SUMMARY

Master's student in Computer Science & Engineering (graduating **Dec 2025**) with a dual background in psychology and expertise spanning **full-stack development** and **applied AI**. Experienced in building **scalable, cloud-native web and mobile applications** from responsive UIs to optimized backend systems, integrating APIs, microservices, and AI-powered features. Skilled in **React, TypeScript, Django, FastAPI, PostgreSQL, GCP**, and modern **DevOps** practices, with hands-on experience designing **multi-agent systems, retrieval-augmented generation (RAG)** pipelines, and production-ready AI integrations that improve usability, performance, and decision-making.

SKILLS

- Programming Languages: TypeScript, JavaScript (ES6+), Python, Java
- Frontend Development: React, React Native (Expo), Next.js, Redux Toolkit, Tailwind CSS, NativeWind, Vite, HTML5, CSS3, Sass, SolidJS, Axios
- Backend & Databases: FastAPI, Django, Express.js, RESTful API & gRPC design, PostgreSQL, Firebase (Firestore), Supabase, Redis, Kafka, RabbitMQ, Microservices Architecture
- Infrastructure & Cloud: Google Cloud Platform (GCP), AWS, Docker, Kubernetes, Terraform, Vercel, Postman
- DevOps & Testing: Git, GitHub Actions, CI/CD pipelines, Pytest, Jest, Jenkins, Unit Testing, Integration Testing, End-to-End (E2E) Testing, Test Automation
- Data & Visualization: Pandas, NumPy, Matplotlib, Seaborn, Plotly, Chart.js, Prometheus, Grafana, Logging & Monitoring
- Generative AI & Tooling: OpenAI API, Gemini API, Claude, Hugging Face Transformers, Ollama, AutoGen, LangChain, LlamaIndex, Chroma, Pinecone, LastMile AI, Weights & Biases (W&B), Cursor

EXPERIENCE

Founding Engineer @ Ripplet | June 2024 – Present

- Developed and deployed Ripplet's marketing site with **Next.js, React, Tailwind CSS**, and **Vercel**, achieving pixel-perfect brand consistency, **responsive design**, and optimized performance with **SEO** best practices
- Integrated **multilingual support** using **i18next**, enhancing accessibility and usability for our diverse user base
- Built responsive dashboards with **Chart.js** to visualize **real-time** client sentiment trends, progress milestones, and engagement levels, helping therapists review sessions more effectively and tailor future treatment plans
- Architected and implemented a **HIPAA-compliant, multi-agent system** with **retrieval-augmented generation (RAG)** pipelines to analyze session transcripts and map client narratives to evidence-based psychological frameworks
- Implemented a secure API layer for session data ingestion and retrieval, optimizing database access with **PostgreSQL** indexing and **Redis** caching, **reducing query latency by 120ms (30%)** in testing compared to unoptimized baselines
- Integrated asynchronous processing with **Kafka** and **RabbitMQ** for large-volume sentiment analysis and contextual resource retrieval, maintaining stable performance under simulated high-load conditions in staging

Full Stack Engineer @ Tech4Good Lab | June 2023 – Jan 2025

- Led a **cross-functional** team of 10 engineers and designers to develop Pathways, an **AI-powered** self-directed learning platform, coordinating backend integration and pilot testing with **over 100 university students**
- Delivered the core application using **Solid.js, Express.js**, and **Firebase**, implementing **server-side rendering, lazy loading**, and **cached API responses**, **reducing average page load times by 35%** in staging tests
- Containerized services with **Docker** and configured **CI/CD pipelines** via **GitHub Actions** for automated build, test, and deployment to staging environments, **cutting manual deployment time from hours to under 10 minutes**
- Developed and maintained **unit and integration tests** for backend services using **Jest** and **Pytest**, **achieving over 85% test coverage** and ensuring **stable performance under simulated high-load scenarios**
- Designed and refined LLM prompts with **different prompting strategies**, validated via **A/B testing**, **reducing hallucinations** and **improving course recommendation relevance by 15%** as measured by completion rates

Backend Developer Intern @ WayOps | July 2021 - August 2021

- Aligned backend deliverables with business goals by collaborating with engineers and stakeholders, ensuring data infrastructure changes met product timelines and operational requirements
- **Increased query performance by 40%** in a production relational database by applying **indexing, query restructuring**, and **strategic caching**, reducing latency and supporting high-throughput analytical workloads
- Enhanced database scalability by **refactoring schema** to remove redundant fields and reorganize table relationships, **reducing duplication by 30%** and improving write throughput under higher volumes and concurrent access

Coding Instructor @ Code For Fun | Feb 2023 - Feb 2024

- Taught programming to cohorts of 1–300 students (ages 6–18), fostering an interactive and inclusive environment
- Developed adaptive, project-based curriculum with **90%** parent satisfaction, guiding students to build websites, data-driven applications, and automation scripts to strengthen **real-world problem-solving skills**

PROJECTS

Large Language Models (LLMs) are Autonomous Cyber Defenders (ACD) | Python

- Researched LLMs as autonomous cybersecurity agents, co-authoring an **IEEE CAI 2025** paper on **explainability and decision transparency**, presented at the conference and **published on arXiv: 2505.04843**
- Extracted and embedded **500+ action-reason statements** using **OpenAI's Embeddings API**, converting LLM-generated rationales into high-dimensional vectors to enable downstream clustering and pattern analysis
- Applied **unsupervised ML algorithms (K-Means, DBSCAN, PCA)** with **feature standardization and dimensionality reduction**, uncovering 5+ interpretable behavioral clusters in autonomous agent decision-making
- Built a reasoning summarizer with **OpenAI GPT-4o** that transformed clustered behaviors into human-readable defense strategies via advanced prompting, **improving explainability and transparency** for LLM-driven autonomous systems.

Travel Agent | React Native, NativeWind, Redux, AutoGen, FastAPI, PostgreSQL, Redis, GCP

- Designed and built a **full-stack mobile travel app** with a **multi-agent system**, integrating real-time web search, API tool calling, itinerary generation, and a multi-turn critic to deliver personalized and constraint-aware plans
- Built a robust agentic web scraping module using **Perplexica** for search discovery, **Playwright** for dynamic rendering, and **Trafilatura** for clean content extraction, **improving relevant content retrieval accuracy by 35%** over baseline
- Built a **rule-based and semantic-relevance filtering** layer to evaluate **open-domain** scraped content for freshness, factual accuracy, and domain relevance, **increasing usable content by 35%** over baseline
- **Improved travel plan quality beyond a GPT-4o baseline** by integrating a multi-turn critic agent to review filtered web content, search results, and generated itineraries, producing more structured, detailed, and user-aligned plans
- Implemented **Redis-backed** state management for cross-agent memory and caching, deployed on **GCP** with containerized services and managed **PostgreSQL**, ensuring low-latency queries and high availability in production

Yi | React, Next.js, Tailwind CSS, Django, REST API, PostgreSQL

- Built a **B2B SaaS application** for small businesses to unify business management and financial analytics, featuring interactive dashboards and visualizations for real-time performance tracking
- Implemented secure **large-file upload** workflows with **Google Cloud Storage signed URLs** and **resumable upload handling**, providing **real-time progress feedback** and ensuring transactional integrity for financial record verification
- Engineered a **reusable component library** with built-in **theme switching** using **Tailwind CSS** configuration, ensuring consistent UI patterns, improving codebase **maintainability**, and facilitating faster feature development
- Developed **authentication** and **role-based access control (RBAC)** with **Django REST Framework** and **JWT**, enabling differentiated permissions for staff, admins, and guests to securely access relevant data and features
- Integrated **Redux Toolkit** and **RTK Query** for centralized **state management** and efficient backend integration, utilizing **caching** and **request deduplication** to improve performance and support scalable data flows

NoteGrid | React, Next.js, Tailwind CSS, Supabase

- Built a note-taking web application with switchable modes for rich text or code editing, integrated with a calendar that categorizes and color-tags notes for chronological tracking
- Built a **TipTap**-based text editor in **Next.js** with real-time font and style controls powered by **state-driven UI updates**, storing content as structured JSON for consistent **cross-device rendering** and **fast page-load rehydration**
- Integrated **Monaco Editor** for in-browser code editing with real-time execution via the **Piston API**, implementing multi-language syntax highlighting, and custom themes to improve performance and user experience
- Implemented an **AI-powered** note analysis and summarization feature using the **OpenAI APIs** with **advanced prompting** to generate concise summaries and extract actionable insights directly from user notes

EDUCATION

University of California, Santa Cruz

Computer Science and Engineering, Master of Science (M.S.)

Expected Dec 2025

Computer Science and Engineering, Bachelor of Science (B.S.); Psychology, Bachelor of Art (B.A.)

Dec 2022